

HW1

October 15, 2018

```
In [1]: import urllib
import numpy as np
from sklearn.svm import LinearSVC
```

1 Regression

```
In [2]: def parseData(fname):
    for l in urllib.request.urlopen(fname):
        yield eval(l)

data = list(parseData("http://jmcauley.ucsd.edu/cse255/data/beer/beer_50000.json"))

reviews = {}
for d in data:
    reviews[d['review/taste']] = reviews.get(d['review/taste'], 0) + 1

reviews
```

```
Out[2]: {1.0: 211,
1.5: 343,
2.0: 1099,
2.5: 1624,
3.0: 4137,
3.5: 8797,
4.0: 16575,
4.5: 12883,
5.0: 4331}
```

```
In [3]: feature1 = [d['beer/style']=='Hefeweizen' for d in data]
feature2 = [d['beer/ABV'] for d in data]
target = [d['review/taste'] for d in data]

X = np.stack((np.ones(len(data)), feature1, feature2), -1)
Y = np.array(target).reshape(-1, 1)

theta, residuals, rank, s = np.linalg.lstsq(X, Y, rcond=None)
theta
```

```
# Theta0 represents bias, theta1 and theta2 represent how much
# weight 'beer' are a Hefeweizen' and 'beer/ABV' contribute
# to taste prediction.
```

```
Out[3]: array([[ 3.11795084],
               [-0.05637406],
               [ 0.10877902]])
```

```
In [4]: def split_data(X, Y, holdout_ratio, shuffle=False):
        m = X.shape[0]
        m_holdout = int(m * holdout_ratio)

        if shuffle:
            permutation = np.random.permutation(m)
            X = X[permutation, :]
            Y = Y[permutation, :]

        return X[:-m_holdout, :], Y[:-m_holdout, :], X[-m_holdout:, :], Y[-m_holdout:, :]

X_train, y_train, X_test, y_test = split_data(X, Y, 0.5)
theta, residuals, rank, s = np.linalg.lstsq(X_train, y_train, rcond=None)

mse_train = np.mean((np.dot(X_train, theta) - y_train)**2)
mse_test = np.mean((np.dot(X_test, theta) - y_test)**2)

print('Train MSE: %f Test MSE: %f' % (mse_train, mse_test))
```

```
Train MSE: 0.483968 Test MSE: 0.423707
```

```
In [5]: X_train, y_train, X_test, y_test = split_data(X, Y, 0.5, shuffle=True)
        theta, residuals, rank, s = np.linalg.lstsq(X_train, y_train, rcond=None)

mse_train = np.mean((np.dot(X_train, theta) - y_train)**2)
mse_test = np.mean((np.dot(X_test, theta) - y_test)**2)

print('Train MSE: %f Test MSE: %f' % (mse_train, mse_test))

# In previous experiment, first half of data may be harder to predict than the
# second half. Random split may reduce bias in original dataset distribution.
```

```
Train MSE: 0.454771 Test MSE: 0.444600
```

```
In [6]: feature1 = [d['beer/ABV'] if d['beer/style']=='Hefeweizen' else 0 for d in data]
        feature2 = [d['beer/ABV'] if d['beer/style']!='Hefeweizen' else 0 for d in data]
        X = np.stack((np.ones(len(data)), feature1, feature2), -1)
```

```

X_train, y_train, X_test, y_test = split_data(X, Y, 0.5)
theta, residuals, rank, s = np.linalg.lstsq(X_train, y_train, rcond=None)

mse_train = np.mean((np.dot(X_train, theta) - y_train)**2)
mse_test = np.mean((np.dot(X_test, theta) - y_test)**2)

print('Train MSE: %f Test MSE: %f' % (mse_train, mse_test))

```

Train MSE: 0.483967 Test MSE: 0.423697

In [7]: *# The features are different. In Q2-4, 'beer/style is Hefeweizen' is showed
explicitly, while in Q5 it is implied from both features.*

2 Classification

```

In [8]: feature1 = [d['review/taste'] for d in data]
        feature2 = [d['review/appearance'] for d in data]
        feature3 = [d['review/aroma'] for d in data]
        feature4 = [d['review/palate'] for d in data]
        feature5 = [d['review/overall'] for d in data]
        target = [d['beer/style'] == 'Hefeweizen' for d in data]

X = np.stack((np.ones(len(data)), feature1, feature2, feature3, feature4, feature5), -1)
Y = np.array(target, dtype=np.int).reshape(-1, 1)

X_train, y_train, X_test, y_test = split_data(X, Y, 0.5, shuffle=True)

clf = LinearSVC(C=1000)
clf.fit(X_train, y_train.reshape(-1))

train_accuracy = clf.score(X_train, y_train)
test_accuracy = clf.score(X_test, y_test)
print('Train accuracy: %f Test accuracy: %f' % (train_accuracy, test_accuracy))

```

Train accuracy: 0.987000 Test accuracy: 0.988280

/home/fanjin/.local/lib/python3.5/site-packages/sklearn/svm/base.py:922: ConvergenceWarning: Lib
"the number of iterations.", ConvergenceWarning)

```

In [9]: text_hef = [d['review/text'] for d in data if d['beer/style'] == 'Hefeweizen']
        text_non_hef = [d['review/text'] for d in data if d['beer/style'] != 'Hefeweizen']

import re
import nltk
from nltk.corpus import stopwords

```

```

# nltk.download('stopwords', download_dir='.')
nltk.data.path.append(".");

def createCorpus(text):
    corpus = {}
    for t in text:
        words = re.findall(r'[\s()!,.?":;0-9]+', t.lower())
        for w in set(words):
            corpus[w] = corpus.get(w, 0) + 1

    remove = [k for k in corpus.keys() if k in stopwords.words('english')]
    for k in remove: del corpus[k]
    return corpus

hef_corpus = createCorpus(text_hef)
nonhef_corpus = createCorpus(text_non_hef)

hef_top_words = [w[0] for w in sorted(hef_corpus.items(), key=lambda x: x[1],
                                     reverse=True)[:50]]
nonhef_top_words = [w[0] for w in sorted(nonhef_corpus.items(),
                                         key=lambda x: x[1], reverse=True)[:300]]

token = [w for w in hef_top_words if w not in nonhef_top_words]
token

# Extract top 50 most frequent occurrence words in
# Hefeweizen review text (exclude stopwords)
# minus top 300 most frequent occurrence words in non-Hefeweizen review text.

```

```
Out[9]: ['banana', 'wheat', 'clove', 'hefe', 'lemon', 'hefeweizen']
```

```

In [10]: def hasToken(text):
        words = re.findall(r'[\s()!,.?":;0-9]+', text.lower())
        cnt = 0
        for w in set(words):
            if w in token:
                cnt += 1
        return cnt

feature6 = [hasToken(d['review/text']) for d in data]

# Feature6 counts how many token words occur in review texts.

In [12]: X = np.stack((np.ones(len(data)), feature1, feature2, feature3,
                        feature4, feature5, feature6), -1)
Y = np.array(target, dtype=np.int).reshape(-1, 1)

```

```
X_train, y_train, X_test, y_test = split_data(X, Y, 0.5, shuffle=True)

clf = LinearSVC(C=1000)
clf.fit(X_train, y_train.reshape(-1))

train_accuracy = clf.score(X_train, y_train)
test_accuracy = clf.score(X_test, y_test)
print('Train accuracy: %f Test accuracy: %f' % (train_accuracy, test_accuracy))
```

```
Train accuracy: 0.991480 Test accuracy: 0.991000
```

```
/home/fanjin/.local/lib/python3.5/site-packages/sklearn/svm/base.py:922: ConvergenceWarning: Lib  
"the number of iterations.", ConvergenceWarning)
```

```
In [ ]:
```