

HW2

October 29, 2018

```
In [1]: import numpy as np
import urllib
import networkx as nx
import matplotlib.pyplot as plt
from homework2_starter import *

%load_ext autoreload
%autoreload 2
```

1 Classifier evaluation

```
In [2]: data = list(parseData("http://jmcauley.ucsd.edu/cse190/data/beer/beer_50000.json"))
X = [feature(d) for d in data]
y = [d['beer/ABV'] >= 6.5 for d in data]

X = np.array(X)
y = np.array(y).reshape(-1, 1)
```

```
In [3]: #####
# 1                                           #
#####

def split_data(X, Y, val_ratio, test_ratio, shuffle=False):
    m = X.shape[0]

    m_train = int(m * (1 - val_ratio - test_ratio))
    m_val = int(m * val_ratio) + m_train
    m_test = int(m * test_ratio) + m_val

    if shuffle:
        permutation = np.random.permutation(m)
        X = X[permutation, :]
        Y = Y[permutation, :]

    return (X[:m_train, :], Y[:m_train, :], X[m_train:m_val, :],
            Y[m_train:m_val, :], X[m_val:m_test, :], Y[m_val:m_test, :])
```

```

def evaluation(theta, X, y):
    scores = [inner(theta,x) for x in X]
    predictions = [s > 0 for s in scores]
    correct = [(a==b) for (a,b) in zip(predictions,y)]
    acc = sum(correct) * 1.0 / len(correct)
    return acc

X_train, y_train, X_val, y_val, X_test, y_test = split_data(X, y, 1/3, 1/3, True)

lam = 1.
theta = train(lam, X, X_train, y_train)
acc = evaluation(theta, X_val, y_val)
print("Validation set accuracy=%f" % acc)
acc = evaluation(theta, X_test, y_test)
print("Test set accuracy=%f" % acc)

```

Validation set accuracy=0.717749
Test set accuracy=0.718229

```

In [4]: #####
# 2 #
#####
def confusion_mat(theta, X, y):
    scores = [inner(theta,x) for x in X]
    predictions = [s > 0 for s in scores]

    tp = 0
    fp = 0
    fn = 0
    tn = 0

    for (a, b) in zip(predictions, y):
        if a == b:
            if a == True: tp += 1
            else: tn += 1
        else:
            if a == True: fp += 1
            else: fn += 1
    return tp, fp, fn, tn

tp, fp, fn, tn = confusion_mat(theta, X_test, y_test)
print('Positive:%d\tNegative:%d\nTP:%d\tTN:%d\tFP:%d\tFN:%d\t' %
      (tp+fp, fn+tn, tp, tn, fp, fn))

```

Positive:12324 Negative:4342
TP:8997 TN:2973 FP:3327 :FN:1369

```

In [5]: #####
# 3                                           #
#####

# To assign more importance to FP, I increase the cost
# (negative loglikelihood) by 10 when FP prediction occurs.

def f2(theta, X, y, lam):
    loglikelihood = 0
    for i in range(len(X)):
        logit = inner(X[i], theta)
        t = log(1 + exp(-logit))
        if not y[i]:
            t += logit

        # when FP occurs, increase loglikelihood by 10
        if y[i] == False and logit > 0:
            t *= 10
        loglikelihood -= t

    for k in range(len(theta)):
        loglikelihood -= lam * theta[k]*theta[k]

    return -loglikelihood

def fprime2(theta, X, y, lam):
    dl = [0]*len(theta)
    for i in range(len(X)):
        logit = inner(X[i], theta)
        for k in range(len(theta)):
            t = X[i][k] * (1 - sigmoid(logit))
            if not y[i]:
                t -= X[i][k]

            # modify derivative accordingly
            if y[i] == False and logit > 0:
                t *= 10
            dl[k] += t

    for k in range(len(theta)):
        dl[k] -= lam*2*theta[k]
    return numpy.array([-x for x in dl])

def train2(lam, X, X_train, y_train):
    theta,_,_ = scipy.optimize.fmin_l_bfgs_b(f2, [0]*len(X[0]),
                                             fprime2, pgtol = 10,
                                             args = (X_train, y_train, lam))

    return theta

```

```

lam = 1.
theta = train2(lam, X, X_train, y_train)
tp, fp, fn, tn = confusion_mat(theta, X_test, y_test)
print('Positive:%d\tNegative:%d\nTP:%d\tTN:%d\tFP:%d\tFN:%d\t' %
      (tp+fp, fn+tn, tp, tn, fp, fn))

```

```

Positive:0          Negative:16666
TP:0          TN:6300          FP:0          :FN:10366

```

```

In [6]: #####
# 4                                           #
#####
lambdas = [0, 0.01, 0.1, 1, 100]
best_lam = None
best_acc = {}

for lam in lambdas:
    theta = train(lam, X, X_train, y_train)
    val_acc = evaluation(theta, X_val, y_val)
    if best_lam is None or best_acc['val'] < val_acc:
        best_lam = lam
        best_acc['train'] = evaluation(theta, X_train, y_train)
        best_acc['val'] = val_acc
        best_acc['test'] = evaluation(theta, X_test, y_test)

print('Best lambda = %f' % best_lam)
print('Train acc=%f \tValidation acc=%f \tTest acc=%f' %
      (best_acc['train'], best_acc['val'], best_acc['test']))

```

```

Best lambda = 0.010000
Train acc=0.720929          Validation acc=0.718409          Test acc=0.718829

```

2 Community Detection

```

In [7]: edges = set()
nodes = set()
for edge in urllib.request.urlopen("http://jmcauley.ucsd.edu/cse255/data/facebook/egonet
    x,y = edge.split()
    x,y = int(x),int(y)
    edges.add((x,y))
    edges.add((y,x))
    nodes.add(x)
    nodes.add(y)

```

```

In [8]: #####
# 5                                           #

```

```
#####

def dfs(node):
    if node in visited:
        return
    comp.add(node)
    visited.add(node)
    for edge in edges:
        if node == edge[0]:
            dfs(edge[1])

visited = set()
comps = []
comps_len = []

for node in nodes:
    comp = set()
    dfs(node)
    if len(comp) > 0:
        comps.append(comp)
        comps_len.append(len(comp))

print('largest connected component contains %d nodes' % max(comps_len))
```

largest connected component contains 40 nodes

```
In [9]: #####
# 6                                           #
#####

def cut_cost(s1, s2):
    cut = 0
    for edge in edges:
        x, y = edge
        if x in s1 and y in s2:
            cut += 1

    cost = 0.5 * (cut/len(s1) + cut/len(s2))
    return cost

comp = comps[comps_len.index(max(comps_len))]
comp = list(comp)
comp = sorted(comp)
s1, s2 = comp[:len(comp)//2], comp[len(comp)//2:]

nodes = [n for n in nodes if n in comp]
edges = [e for e in edges if e[0] in nodes or e[1] in nodes]
```

```

nodes = set(nodes)
edges = set(edges)

cost = cut_cost(s1, s2)
print('normalized-cut cost of the 50/50 split: %f' % cost)

```

normalized-cut cost of the 50/50 split: 4.600000

```

In [10]: #####
# 7                                           #
#####

def greedy(s1, s2, compute_cost):
    node_cost = []

    for node in s1:
        t1 = set(s1)
        t2 = set(s2)
        t1.remove(node)
        t2.add(node)
        node_cost.append((node, compute_cost(t1, t2)))

    for node in s2:
        t1 = set(s1)
        t2 = set(s2)
        t1.add(node)
        t2.remove(node)
        node_cost.append((node, compute_cost(t1, t2)))

    return node_cost

In [11]: comp = comps[comps_len.index(max(comps_len))]
comp = list(comp)
comp = sorted(comp)
s1, s2 = set(comp[:len(comp)//2]), set(comp[len(comp)//2:])

cur_cost = cut_cost(s1, s2)

while(True):
    node_cost = greedy(s1, s2, cut_cost)
    node_cost = sorted(node_cost, key=lambda x:(x[1], x[0]))
    if node_cost[0][1] >= cur_cost: break

    cur_cost = node_cost[0][1]
    node = node_cost[0][0]
    if node in s1:
        s1.remove(node)

```

```

        s2.add(node)
    else:
        s1.add(node)
        s2.remove(node)

    print('Elements in split 1:\n' + str(s1))
    print('Elements in split 2:\n' + str(s2))
    print('normalized-cut cost: %f' % cur_cost)

```

```

Elements in split 1:
{769, 772, 774, 798, 800, 803, 805, 810, 811, 819, 823, 697, 828, 830, 703, 708, 840, 713, 719,
Elements in split 2:
{864, 804, 876, 893, 878, 882, 884, 888, 886, 729, 825, 889, 861, 863}
normalized-cut cost: 0.879121

```

```

In [12]: #####
        # 8                                     #
        #####

```

```

def modularity(s):
    e = 0
    a = 0
    for edge in edges:
        x, y = edge
        if x in s and y in s:
            e += 1
        if x in s or y in s:
            a += 1
    e /= len(edges)
    a /= len(edges)

    mod = e - a**2
    return mod

def mod_cost(s1, s2):
    return modularity(s1) + modularity(s2)

```

```

In [13]: comp = comps[comps_len.index(max(comps_len))]
        comp = list(comp)
        comp = sorted(comp)
        s1, s2 = set(comp[:len(comp)//2]), set(comp[len(comp)//2:])

        cur_cost = mod_cost(s1, s2)

        while(True):
            node_cost = greedy(s1, s2, mod_cost)
            node_cost = sorted(node_cost, key=lambda x:(x[1], x[0]), reverse=True)

```

```

    if node_cost[0][1] <= cur_cost: break

    cur_cost = node_cost[0][1]
    node = node_cost[0][0]
    if node in s1:
        s1.remove(node)
        s2.add(node)
    else:
        s1.add(node)
        s2.remove(node)

    print('Elements in split 1:\n' + str(s1))
    print('Elements in split 2:\n' + str(s2))
    print('modularity: %f' % cur_cost)

```

```

Elements in split 1:
{769, 772, 774, 798, 800, 803, 805, 810, 811, 819, 823, 697, 828, 830, 703, 708, 840, 713, 719,
Elements in split 2:
{864, 804, 876, 893, 878, 882, 884, 888, 886, 729, 856, 825, 889, 861, 863}
modularity: 0.233368

```

In []: